

Programming Assignment 1: “Frontend”

1 Overview

Programming assignments 1–3 will direct you to design and build a compiler for Cool. Each assignment will cover one component of the compiler: frontend (lexical and syntactic analysis), semantic analysis, and backend (code generation). Each assignment will ultimately result in a working component, which can interface with other the components. You will implement the compiler in Java.

We strongly recommend that you work in teams of 2 or 3 students. You must work in the same team on all three assignments A1–A3. Please sign up for your team on QM+.

Documentation for the Cool programming language and all the tools needed for building a compiler will be made available on the QM+ page. This includes a manual for `Antlr4`, as well as the manual for the spim simulator.

There is a lot of information in this handout, and you need to know most of it to write a working frontend. *Please read the handout thoroughly.*

1.1 Frontend

For this assignment, you are to write a frontend. It consists of a lexical analyzer (also called a scanner or lexer) and a syntactic analyzer (also called a parser). You will be using an *analyzer generator* called `Antlr4`.

The output of your frontend will be an abstract syntax tree (AST). You will construct this AST using a Listener or Visitor automatically generated by the analyzer generator.

You must make some provision for graceful termination if a fatal error occurs. Do not leave exceptions uncaught.

1.2 Lexical Analysis

For the lexical analyzer you will describe the set of tokens for Cool in the `Antlr4` input format, and the analyzer generator will generate the actual Java code for recognizing tokens in Cool programs.

You should follow the specification of the lexical structure of Cool given in the Cool Reference Manual. Your scanner should be robust—it should work for any conceivable input. For example, you must handle errors such as EOF occurring in the middle of a string or comment, as well as string constants that are too long. These are just some of the errors that can occur; see the Cool Reference Manual for the rest.

Note that if the lexical specification is incomplete (some input has no regular expression that matches), then the generated scanner does undesirable things. *Make sure your specification is complete.*

For class identifiers, object identifiers, integers, and strings, the semantic value should be of type `Symbol`. For boolean constants, the semantic value is of type `java.lang.Boolean`. When a lexical error is encountered, the semantic value is the error message string. The lexemes for the other tokens do not carry any interesting information.

There is an issue in deciding how to handle the special identifiers for the basic classes (`Object`, `Int`, `Bool`, `String`), `SELF_TYPE`, and `self`. However, this issue doesn’t actually come up until later phases of the compiler—the scanner should treat the special identifiers exactly like any other identifier.

Do *not* test whether integer literals fit within the representation specified in the Cool Reference Manual—simply create a `Symbol` with the entire literal text as its content, regardless of its length.

Your lexer should maintain a variable that indicates which line in the source text is currently being scanned.

This feature aids the parser in printing useful error messages. This variable is generated and maintained automatically by Antlr4.

Programs tend to have many occurrences of the same lexeme. For example, an identifier is generally referred to more than once in a program. To save space and time, a common compiler practice is to store lexemes in a *string table*. We provide a string table implementation in Java. For the moment, you only need to know that the type of string table entries is `Symbol`.

1.2.1 Strings

Your scanner should convert escape characters in string constants to their correct values. For example, if the programmer typed these eight characters:

" a b \ n c d "

your scanner would return the token `STR_CONST` whose semantic value is these 5 characters:

a b \n c d

where `\n` represents the literal ASCII character for newline.

Following the specification on page 15 of the Cool Reference Manual, you must return an error for a string containing the literal null character. However, a sequence of the two characters

\ 0

is allowed and should be converted to the one character

0 .

1.3 Syntactic Analysis

You will need to refer to the syntactic structure of Cool, found in the Cool Reference Manual.

Your frontend should build an AST using the `ast` and `ast.visitor` packages provided. You will need this data structure for this and future assignments. The root (and only the root) of the AST should be of type `ProgramNode`. For programs that are parsed successfully, the output of the `parser` is a listing of the AST.

Your parser need only work for programs contained in a single file—don’t worry about compiling multiple files.

For programs that have errors, the output is the error messages of the parser. We have supplied an error reporting routine that prints error messages in a standard format; please do not modify it.

Since the `mycoolc` compiler uses pipes to communicate from one stage to the next, any extraneous characters produced by the parser can cause errors; in particular, the semantic analyzer may not be able to read the AST your parser produces.

You may use precedence declarations, but only for expressions. Do not use precedence declarations blindly (i.e., do not respond to failure of Antlr to generate an analyzer from your grammar by adding precedence rules until it goes away).

The Cool `let` construct introduces an ambiguity into the language (try to construct an example if you are not convinced). The Cool Reference Manual resolves the ambiguity by specifying that a `let` expression extends as far to the right as possible.

1.4 Error Handling

The purpose of error handling is to permit the parser to continue after some anticipated error. It is not a panacea and the parser may become completely confused. You can use the error handling capabilities of Antlr4. Refer to the Antlr4 documentation for details.

There are several requirements for reporting and recovering from lexical errors:

- When an invalid character (one that can't begin any token) is encountered, a string containing just that character should be returned as the error string. Resume lexing at the following character.
- If a string contains an unescaped newline, report that error as "**Unterminated string constant**" and resume lexing at the beginning of the next line—we assume the programmer simply forgot the close-quote.
- When a string is too long, report the error as "**String constant too long**" in the error string in the **ERROR** token. If the string contains invalid characters (i.e., the null character), report this as "**String contains null character**" or "**String contains escaped null character**". In either case, lexing should resume after the end of the string. The end of the string is defined as either
 1. the beginning of the next line if an unescaped newline occurs after these errors are encountered;
or
 2. after the closing " otherwise.
- If a comment remains open when EOF is encountered, report this error with the message "**EOF in comment**". Do *not* tokenize the comment's contents simply because the terminator is missing. Similarly for strings, if an EOF is encountered before the close-quote, report this error as "**EOF in string constant**".
- If you see “***)**” outside a comment, report this error as "**Unmatched *)**", rather than tokenizing it as ***** and **)**.

Do not be overly concerned about the the line numbers that appear in the error messages your parser generates. If your parser is working correctly, the line number will generally be the line where the error occurred. For erroneous constructs broken across multiple lines, the line number will probably be the last line of the construct.

In your **README**, describe which errors you attempt to catch. Your test file **bad.cl** should have some instances that illustrate the errors from which your parser can recover.

1.5 Testing

The first way to test your scanner is to generate sample inputs and run them using **myfrontend -x**, which prints out the line number and the lexeme of every token recognized by your scanner.

To test the parser, you will need a working scanner. Don't automatically assume that the scanner is bug free—latent bugs in the scanner may cause mysterious problems in the parser.

You can run your frontend using **myfrontend**. When you think your frontend is working, try running **mycoold** to invoke your frontend together with the other compiler components provided. This will be a complete Cool compiler that you can try on any test program.

You should test this compiler on both good and bad inputs to see if everything is working. Remember, bugs in your parser may manifest themselves elsewhere. To run the testsuite, use **testme lexer parser**. Make sure to inspect the logs and detailed outputs generated by the script. Do not rely only on PASS/FAIL counts.

1.6 String Tables

All compilers manage large numbers of strings such as program identifiers, numerical constants, and string constants. Often, many of these strings are the same. For example, each identifier typically occurs many

times in a program. To ensure that string constants are stored compactly and manipulated efficiently, a specialized data structure, the *string table*, is employed.

A string table is a lookup table that maintains a single copy of each string. The `StringTable` class provides methods for inserting and querying string tables in a variety of ways. It is implemented using a `HashMap`. Each entry in the string table stores a string and an integer index unique to the string.

An important point about the structure of the Cool compiler is that there are actually three distinct string tables: one for string constants (`stringtable`), one for integer constants (`inttable`), and one for identifiers (`idtable`). The code generator must distinguish integer constants and string constants from each other and from identifiers, because special code is produced for each string constant and each integer constant in the program. Having three distinct string tables makes this distinction easy. Throughout the rest of the compiler (except parts of the code generator), a pointer to an entry in a string table is called a `Symbol`, irrespective of whether the symbol represents an integer, string, or identifier.

Because string tables store only one copy of each string, comparing whether two entries `x` and `y` represent the same string is simple. Note that it does not make sense to compare entries from different string tables, as these are guaranteed to be different even if the strings are the same. The `TreeConstants` class uses string tables.

1.7 Abstract Syntax Trees

After lexical analysis and parsing, a Cool program is represented internally by the Cool compiler as an abstract syntax tree (AST).

The template code for this assignment includes the definition of the AST data type, in the package `ast`. The AST data type provides, for each kind of Cool construct, a class for representing constructs of that kind. There is a class for `let` expressions, another class for `+` expressions, and so on. Objects of these classes are nodes in Cool abstract syntax trees. For example, an expression `e1 + e2` is represented by a `+` expression object, which has two subtrees: one for the tree representing the expression `e1` and one for the tree representing the expression `e2`.

The template code also provides a convenient way to implement AST traversals using the Visitor Design Pattern, in the package `ast/visitors`. For example, the `DumpVisitor` pretty-prints an AST. You will implement other visitors in A2 and A3.

1.7.1 AST Class Hierarchy

All AST classes are derived from the class `TreeNode`, which has one field for the line number in which the construct corresponding to the AST node appeared in the source file. The line number is used by the compiler for producing descriptive error messages.

Each class definition in the `ast` package comes with a number of fields and getter methods for these fields. Fields are only visible to methods of that class or derived classes. All classes define `accept` methods for the Visitor Design Pattern.

Each class corresponds to a portion of the Cool grammar. The fields of each class correspond to non-terminals and terminals that appear in productions in the Cool syntax specification in the manual. This correspondence between the AST data type and the Cool program syntax should make clear how to use the AST classes.

A sample class definition is

```
public class ClassNode extends TreeNode {  
    protected Symbol name;  
    protected Symbol parent;
```

```
protected List<FeatureNode> features = new LinkedList<FeatureNode>();  
protected Symbol filename;  
...  
}
```

An object of `ClassNode` is a node with four children: a `Symbol` (a type identifier) for the class name, a `Symbol` (another type identifier) for the parent class, a list of `FeatureNodes`, and a `Symbol` for the filename in which the class definition occurs.

1.7.2 AST Classes

This section briefly describes each class and its role in the compiler. It may be helpful to read this section in conjunction with the code.

- **ProgramNode** This class is for the root of the AST. At the end of parsing, the root holds the final list of classes. The only needed use of this class is already in the template.
- **ClassNode** This class is for AST nodes for Cool classes. See the example above.
- **MethodNode** This is one of the two classes derived from `FeatureNode`. Use this class to build AST nodes for methods. It holds method name, return type, list of formal parameters, and the body of the method.
- **AttributeNode** This is one of the two classes derived from `FeatureNode`. Use this class to build AST nodes for attributes. The `init` field is for the expression that is the optional initialization.
- **FormalNode** This class is for formal parameters in method definitions. The fields are name and declared type of the formal parameter.
- **BranchNode** Use this class to build an AST node for each branch of a `case` expression. A branch has the form

```
name : typeid => expr;
```

which corresponds to the field names of `BranchNode` in the obvious way.

- **AssignNode** This is the class for assignment expressions.
- **StaticDispatchNode** and **DispatchNode** There are two different kinds of dispatch in Cool and they have distinct classes. See the Cool Reference Manual for a discussion of static vs. dynamic dispatch. The Cool syntax has a shorthand for dispatch that omits the `self` parameter. Don’t use `NoExpressionNode` in place of `self`; you need to fill in the symbol for `self` for the rest of the compiler to work correctly.
- **CondNode** This is the class for `if-then-else-fi` expressions.
- **LoopNode** This is the class for `while-loop-pool` expressions.
- **CaseNode** This class builds an AST for a `case` expression. It holds a list of case branches. See `BranchNode` above.
- **BlockNode** This is the class for `{ ... }` block expressions.
- **LetNode** This is the class for `let` expressions. Note that the `LetNode` only allows one identifier. When parsing a `let` expression with multiple identifiers, it should be transformed into nested `lets` with single identifiers, as described in the semantics for `let` in the Cool Reference Manual.

- **PlusNode** This is the class for + expressions.
- **SubNode** This is the class for - expressions.
- **MulNode** This is the class for * expressions.
- **DivideNode** This is the class for / expressions.
- **NegNode** This is the class for ~ expressions.
- **LINode** This is the class for < expressions.
- **EqNode** This is the class for = expressions.
- **LEqNode** This is the class for <= expressions.
- **CompNode** This is the class for not expressions.
- **IntConstNode** This is the class for integer constants.
- **BoolConstNode** This is the class for boolean constants.
- **StringConstNode** This is the class for string constants.
- **NewNode** This is the class for new expressions.
- **IsVoidNode** This is the class for isvoid expressions.
- **ObjectNode** This class is for expressions that are just object identifiers. Object identifiers are used in many places in the syntax, but there is only one production for object identifiers as expressions.
- **NoExpressionNode** Use this class for optional initialization expressions.

2 Getting and Submitting the Assignment

2.1 Initial Group Setup

1. Form a group (recommended and maximum size 3). Tell Raymond your group members and you will be assigned a group number.
2. **One person** should do the following.

- Create a new **private** QMUL Github repository from the *template repository* and add only the other group members. Name the repository `ECS652U-cw-group<number>`. Also add Raymond (`r.hu@qmul.ac.uk`) to your repository.
- Set up a fresh local copy of the template repository and link it to the group QMUL Github repository as in **A0**.
- Before modifying anything in the code project, create a Git branch called `frontend` and link it to the QMUL Github repository as follows. We assume the repository directory is `$COOL_HOME`.

```
cd $COOL_HOME  
git checkout -b frontend  
git push -u origin frontend
```

Check this branch is visible and correct on the QMUL Github Web interface before proceeding further.

3. **The other group members** should then obtain a local instance of this repository using the `git clone` command; see the Github Web interface for example instructions. **After cloning the repository, be sure to switch to the frontend branch** and link it to QMUL Github.

```
cd $COOL_HOME  
git checkout frontend  
git push -u origin frontend
```

You are all now set up to work on your COOL compiler. You have local repositories and a QMUL Github repository that you should use to save your work, collaborate with your team members, and develop your final submission for this assignment. **All work on this assignment by all team members should be done and submitted in the frontend branch**.

4. Before proceeding further, all group members must synchronize to confirm every member knows how to pull team updates from the QMUL Github server, commit their own changes and push to the server (e.g., the `git pull`, `git commit -a -m "...commit message..."`, and `git push` commands). You must agree a “protocol” for determining who is working on what part of the project at any one time, to prevent concurrent conflicts.

In *every* group assignment, all team members are expected to contribute a substantial amount of concrete code to the project, and to understand the entirity of the code (not just the part they implemented). This will be *individually* verified as part of the assessment.

2.2 A1 Instructions and Submission

1. To compile and run the template code for this assignment, run

```
cd $COOL_HOME/assignments/pa1  
buildme frontend  
.myfrontend good.cl
```

Try it straight away—it should “work”, and print some error messages about `token recognition error`, because the lexer is not implemented yet.

2. Check the instructions in the `assignments/pa1/README` file. It has many more details, including testing.
3. During your work, you can compare your output against that of the reference compiler by running `./reffrontend` as opposed to `./myfrontend`. Make your output to match the reference version to pass all the `testme` tests.
4. The files that you will need to modify are:
 - `CoolLexer.g4`
 - `CoolParser.g4`
 - `ASTBuilder.java`

You may also modify the following files, if you wish:

- `TreeConstants.java`
- `CoolErrorListener.java`
- `CoolErrorStrategy.java`

You should NOT edit any other files, nor add any other files that are required to compile your project. In fact, if you modify any other files, you may find it impossible to complete the assignment.

5. You will submit a zip archive of your entire `ECS652U-cw-group<number>` code repository, including the (possibly “hidden”) `.git` directory. Failure to include the `.git` directory will count as an incomplete submission.

To reiterate: we will mark your project based on the latest commit on the `frontend` branch (regardless of the commit message) at the deadline.